




294 visual tokens (14x21 cells)
the tablecloth spans 2 cells



600 visual tokens (20x30 cells)
the tablecloth spans 2 cells



888 visual tokens (24x37 cells)
the tablecloth spans 4 cells



target region as encoded
19 px wide, inside a 679x452 image



target region as encoded
26 px wide, inside a 960x640 image



target region as encoded
32 px wide, inside a 1176x784 image



300 tokens: the tablecloth spans 2 cells — sub-token evidence, the model cannot read the colour. 600 (the equal-compute bar): still 2 cells.

900 tokens: 4 cells, readable — but 900 x 28 layers is 50% over budget. AVR buys it by dropping the last 11 layers' tokens (900x17 + 90x11 = 16,290 = 97% of the bar); DWA instead spends the same 600 tokens as 300 localise + 300 crop at 16x density.