




allowance: $N(b+1) + N_{\text{keep}}(L-b-1) = 900 \times 17 + 90 \times 11 = 16,290$ visual token-layers $\approx 97\%$ of the equal-compute reference ($600 \times 28 = 16,800$)

the saved late-layer compute buys the higher-resolution input: 900 tokens instead of 600, with direct access kept only where it is causally live