


SINGLE-INSTANCE — direct_attributes/71
"What is the color of the paraglider?"
disp = 0.038 > median tau = 0.027 -> CROP




arm outcomes on this item

bar (600, no crop)	X
block mean (published read-out)	X
DWA (ridge crop W=0.25)	OK
AVR (900 + prune)	OK

green box = ground truth; dashed box = where DWA would crop; heat = block-mean attention

CROSS-INSTANCE — relative_position/174
"Is the slide on the left or right side of the life buoy?"
disp = 0.017 < median tau = 0.027 -> FULL-RES (AVR)



arm outcomes on this item

bar (600, no crop)	OK
block mean (published read-out)	X
DWA (ridge crop W=0.25)	X
AVR (900 + prune)	OK

green box = ground truth; dashed box = where DWA would crop; heat = block-mean attention

The policy reads only the concentration of the model's own attention: a tight map (disp above the median) means one region holds the evidence, so crop it at 16x density; a dispersed map means the evidence is spread over the scene, so keep the whole image and buy resolution instead (the crop is what the dashed box would have shown).