



TRAINING (once, ~50 boxed examples)



INFERENCE (per question; one localise pass + one answer pass)



COST: 300 (localise) + 300 (answer) = 600 visual tokens = the equal-compute bar

