

Where you read the attention map decides the method
Qwen3-VL-2B, $n=191$, every arm at ~ 600 visual tokens

