



(a) VLM internals — where the map comes from



(b) DWA — depth-weighted read-out and crop



cost: 300 (localise) + 300 (answer) = 600 visual tokens = the equal-compute bar

the ridge runs on the full depth: the causally inert late layers receive small or negative weight, the question-conditioned band is up-weighted