


VLM internals + DWA — how the ridge read-out fits into the model



cost: 300 (localise) + 300 (answer) = 600 visual tokens = the equal-compute bar

the ridge runs on the full depth: the causally inert late layers receive small or negative weight, the question-conditioned band is up-weighted